

MaCVi

October 25, 2022

Learning resource: Keras Official Website (https://keras.io/examples/vision/oxford_pets_image_s)

0.1 Importing and loading the image sets

```
[3]: import os
      #importing all the functions from function.py
      from function import *
      from tensorflow import keras
      import numpy as np
      from tensorflow.keras.preprocessing.image import load_img
      from IPython.display import Image, display
      import numpy as np
      from tensorflow.keras.preprocessing.image import load_img
      from PIL import ImageOps
      import random

      input_dir = "MaSTr1325_images_512x384/"
      target_dir = "MaSTr1325_masks_512x384/"
      img_size = (512, 384)
      num_classes = 4
      batch_size = 32

      input_img_paths = sorted(
          [
              os.path.join(input_dir, fname)
              for fname in os.listdir(input_dir)
              if fname.endswith(".jpg")
          ]
      )
      target_img_paths = sorted(
          [
              os.path.join(target_dir, fname)
              for fname in os.listdir(target_dir)
              if fname.endswith(".png")
              and not fname.startswith(".")
          ]
      )
```

```
print("Number of samples:", len(input_img_paths))

for input_path, target_path in zip(input_img_paths[:10], target_img_paths[:10]):
    print(input_path, "|", target_path)
```

2022-10-25 22:45:55.835790: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.


```
Number of samples: 1325
MaStr1325_images_512x384/0001.jpg | MaStr1325_masks_512x384/0001m.png
MaStr1325_images_512x384/0002.jpg | MaStr1325_masks_512x384/0002m.png
MaStr1325_images_512x384/0003.jpg | MaStr1325_masks_512x384/0003m.png
MaStr1325_images_512x384/0004.jpg | MaStr1325_masks_512x384/0004m.png
MaStr1325_images_512x384/0005.jpg | MaStr1325_masks_512x384/0005m.png
MaStr1325_images_512x384/0006.jpg | MaStr1325_masks_512x384/0006m.png
MaStr1325_images_512x384/0007.jpg | MaStr1325_masks_512x384/0007m.png
MaStr1325_images_512x384/0008.jpg | MaStr1325_masks_512x384/0008m.png
MaStr1325_images_512x384/0009.jpg | MaStr1325_masks_512x384/0009m.png
MaStr1325_images_512x384/0010.jpg | MaStr1325_masks_512x384/0010m.png
```

0.2 Checking a image and its segmentation mask

```
[4]: # Display input image #7
display(Image(filename=input_img_paths[9]))

# Display auto-contrast version of corresponding target (per-pixel categories)
img = ImageOps.autocontrast(load_img(target_img_paths[9]))
display(img)
```





0.3 Getting the ML model architecture, using UNet

```
[6]: from tensorflow.keras import layers
      # Build model
      model = get_model(img_size, num_classes)
      model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 512, 384, 3)]	0	[]
conv2d_9 (Conv2D)	(None, 256, 192, 4)	112	['input_2[0][0]']
batch_normalization_15 (Batch Normalization)	(None, 256, 192, 4)	16	['conv2d_9[0][0]']

```

ormalization)

activation_15 (Activation)      (None, 256, 192, 4)  0
['batch_normalization_15[0][0]']

activation_16 (Activation)      (None, 256, 192, 4)  0
['activation_15[0][0]']

separable_conv2d_6 (SeparableC  (None, 256, 192, 8)  76
['activation_16[0][0]']
onv2D)

batch_normalization_16 (BatchN  (None, 256, 192, 8)  32
['separable_conv2d_6[0][0]']
ormalization)

activation_17 (Activation)      (None, 256, 192, 8)  0
['batch_normalization_16[0][0]']

separable_conv2d_7 (SeparableC  (None, 256, 192, 8)  144
['activation_17[0][0]']
onv2D)

batch_normalization_17 (BatchN  (None, 256, 192, 8)  32
['separable_conv2d_7[0][0]']
ormalization)

max_pooling2d_3 (MaxPooling2D)  (None, 128, 96, 8)  0
['batch_normalization_17[0][0]']

conv2d_10 (Conv2D)              (None, 128, 96, 8)  40
['activation_15[0][0]']

add_7 (Add)                     (None, 128, 96, 8)  0
['max_pooling2d_3[0][0]',
'conv2d_10[0][0]']

activation_18 (Activation)      (None, 128, 96, 8)  0          ['add_7[0][0]']

separable_conv2d_8 (SeparableC  (None, 128, 96, 16)  216
['activation_18[0][0]']
onv2D)

batch_normalization_18 (BatchN  (None, 128, 96, 16)  64
['separable_conv2d_8[0][0]']
ormalization)

activation_19 (Activation)      (None, 128, 96, 16)  0

```

```

['batch_normalization_18[0][0]']

separable_conv2d_9 (SeparableC (None, 128, 96, 16) 416
['activation_19[0][0]']
onv2D)

batch_normalization_19 (BatchN (None, 128, 96, 16) 64
['separable_conv2d_9[0][0]']
ormalization)

max_pooling2d_4 (MaxPooling2D) (None, 64, 48, 16) 0
['batch_normalization_19[0][0]']

conv2d_11 (Conv2D) (None, 64, 48, 16) 144 ['add_7[0][0]']

add_8 (Add) (None, 64, 48, 16) 0
['max_pooling2d_4[0][0]',
'conv2d_11[0][0]']

activation_20 (Activation) (None, 64, 48, 16) 0 ['add_8[0][0]']

separable_conv2d_10 (Separable (None, 64, 48, 32) 688
['activation_20[0][0]']
Conv2D)

batch_normalization_20 (BatchN (None, 64, 48, 32) 128
['separable_conv2d_10[0][0]']
ormalization)

activation_21 (Activation) (None, 64, 48, 32) 0
['batch_normalization_20[0][0]']

separable_conv2d_11 (Separable (None, 64, 48, 32) 1344
['activation_21[0][0]']
Conv2D)

batch_normalization_21 (BatchN (None, 64, 48, 32) 128
['separable_conv2d_11[0][0]']
ormalization)

max_pooling2d_5 (MaxPooling2D) (None, 32, 24, 32) 0
['batch_normalization_21[0][0]']

conv2d_12 (Conv2D) (None, 32, 24, 32) 544 ['add_8[0][0]']

add_9 (Add) (None, 32, 24, 32) 0
['max_pooling2d_5[0][0]',
'conv2d_12[0][0]']

```

activation_22 (Activation)	(None, 32, 24, 32)	0	['add_9[0][0]']
conv2d_transpose_8 (Conv2DTran ['activation_22[0][0]'] spose)	(None, 32, 24, 32)	9248	
batch_normalization_22 (BatchN ['conv2d_transpose_8[0][0]'] ormalization)	(None, 32, 24, 32)	128	
activation_23 (Activation)	(None, 32, 24, 32)	0	['batch_normalization_22[0][0]']
conv2d_transpose_9 (Conv2DTran ['activation_23[0][0]'] spose)	(None, 32, 24, 32)	9248	
batch_normalization_23 (BatchN ['conv2d_transpose_9[0][0]'] ormalization)	(None, 32, 24, 32)	128	
up_sampling2d_9 (UpSampling2D)	(None, 64, 48, 32)	0	['add_9[0][0]']
up_sampling2d_8 (UpSampling2D) ['batch_normalization_23[0][0]']	(None, 64, 48, 32)	0	
conv2d_13 (Conv2D) ['up_sampling2d_9[0][0]']	(None, 64, 48, 32)	1056	
add_10 (Add) ['up_sampling2d_8[0][0]'], 'conv2d_13[0][0]']	(None, 64, 48, 32)	0	
activation_24 (Activation) ['add_10[0][0]']	(None, 64, 48, 32)	0	
conv2d_transpose_10 (Conv2DTra ['activation_24[0][0]'] nspose)	(None, 64, 48, 16)	4624	
batch_normalization_24 (BatchN ['conv2d_transpose_10[0][0]'] ormalization)	(None, 64, 48, 16)	64	
activation_25 (Activation) ['batch_normalization_24[0][0]']	(None, 64, 48, 16)	0	

```

conv2d_transpose_11 (Conv2DTra (None, 64, 48, 16) 2320
['activation_25[0][0]']
nspose)

batch_normalization_25 (BatchN (None, 64, 48, 16) 64
['conv2d_transpose_11[0][0]']
ormalization)

up_sampling2d_11 (UpSampling2D (None, 128, 96, 32) 0
['add_10[0][0]']
)

up_sampling2d_10 (UpSampling2D (None, 128, 96, 16) 0
['batch_normalization_25[0][0]']
)

conv2d_14 (Conv2D) (None, 128, 96, 16) 528
['up_sampling2d_11[0][0]']

add_11 (Add) (None, 128, 96, 16) 0
['up_sampling2d_10[0][0]',
'conv2d_14[0][0]']

activation_26 (Activation) (None, 128, 96, 16) 0
['add_11[0][0]']

conv2d_transpose_12 (Conv2DTra (None, 128, 96, 8) 1160
['activation_26[0][0]']
nspose)

batch_normalization_26 (BatchN (None, 128, 96, 8) 32
['conv2d_transpose_12[0][0]']
ormalization)

activation_27 (Activation) (None, 128, 96, 8) 0
['batch_normalization_26[0][0]']

conv2d_transpose_13 (Conv2DTra (None, 128, 96, 8) 584
['activation_27[0][0]']
nspose)

batch_normalization_27 (BatchN (None, 128, 96, 8) 32
['conv2d_transpose_13[0][0]']
ormalization)

up_sampling2d_13 (UpSampling2D (None, 256, 192, 16 0
['add_11[0][0]']
)
)

```



```

up_sampling2d_12 (UpSampling2D (None, 256, 192, 8) 0
['batch_normalization_27[0][0]']
)

conv2d_15 (Conv2D) (None, 256, 192, 8) 136
['up_sampling2d_13[0][0]']

add_12 (Add) (None, 256, 192, 8) 0
['up_sampling2d_12[0][0]',
'conv2d_15[0][0]']

activation_28 (Activation) (None, 256, 192, 8) 0
['add_12[0][0]']

conv2d_transpose_14 (Conv2DTra (None, 256, 192, 4) 292
['activation_28[0][0]']
nspose)

batch_normalization_28 (BatchN (None, 256, 192, 4) 16
['conv2d_transpose_14[0][0]']
ormalization)

activation_29 (Activation) (None, 256, 192, 4) 0
['batch_normalization_28[0][0]']

conv2d_transpose_15 (Conv2DTra (None, 256, 192, 4) 148
['activation_29[0][0]']
nspose)

batch_normalization_29 (BatchN (None, 256, 192, 4) 16
['conv2d_transpose_15[0][0]']
ormalization)

up_sampling2d_15 (UpSampling2D (None, 512, 384, 8) 0
['add_12[0][0]']
)

up_sampling2d_14 (UpSampling2D (None, 512, 384, 4) 0
['batch_normalization_29[0][0]']
)

conv2d_16 (Conv2D) (None, 512, 384, 4) 36
['up_sampling2d_15[0][0]']

add_13 (Add) (None, 512, 384, 4) 0
['up_sampling2d_14[0][0]',
'conv2d_16[0][0]']

```

```
conv2d_17 (Conv2D)          (None, 512, 384, 4) 148
['add_13[0][0]']
```

```
=====
=====
Total params: 34,196
Trainable params: 33,724
Non-trainable params: 472
-----
-----
```

0.4 Splitting the dataset into Training set, Cross Validation and Test set

```
[7]: train_samples = int (1325 * 0.7)
      val_samples = train_samples + int (1325 * 0.2)

      random.Random(1337).shuffle(input_img_paths)
      random.Random(1337).shuffle(target_img_paths)
      train_input_img_paths = input_img_paths[:train_samples]
      train_target_img_paths = target_img_paths[:train_samples]
      val_input_img_paths = input_img_paths[train_samples:val_samples]
      val_target_img_paths = target_img_paths[train_samples:val_samples]
      test_input_img_paths = input_img_paths[val_samples:]
      test_target_img_paths = target_img_paths[val_samples:]

      train_gen = SeaKing(
          batch_size, img_size, train_input_img_paths, train_target_img_paths
      )
      val_gen = SeaKing(batch_size, img_size, val_input_img_paths,
          ↪val_target_img_paths)
```

0.5 Training the ML model

```
[6]: model.compile(optimizer="adam", loss="sparse_categorical_crossentropy")

      callbacks = [
          keras.callbacks.ModelCheckpoint("SeaKing.h5", save_best_only=True)
      ]

      # Train the model, doing validation at the end of each epoch.
      epochs = 40
      history = model.fit(train_gen, epochs=epochs, validation_data=val_gen,
          ↪callbacks=callbacks)
```

Epoch 1/40

28/28 [=====] - 110s 4s/step - loss: 0.5022 - val_loss:

```

1.1625
Epoch 2/40
28/28 [=====] - 107s 4s/step - loss: 0.2243 - val_loss:
1.7156
Epoch 3/40
28/28 [=====] - 103s 4s/step - loss: 0.1855 - val_loss:
2.2282
Epoch 4/40
28/28 [=====] - 101s 4s/step - loss: 0.1612 - val_loss:
2.5915
Epoch 5/40
28/28 [=====] - 105s 4s/step - loss: 0.1467 - val_loss:
2.8407
Epoch 6/40
28/28 [=====] - 108s 4s/step - loss: 0.1369 - val_loss:
2.5838
Epoch 7/40
28/28 [=====] - 102s 4s/step - loss: 0.1286 - val_loss:
2.7267
Epoch 8/40
28/28 [=====] - 103s 4s/step - loss: 0.1220 - val_loss:
2.5417
Epoch 9/40
28/28 [=====] - 113s 4s/step - loss: 0.1166 - val_loss:
2.4312
Epoch 10/40
28/28 [=====] - 104s 4s/step - loss: 0.1109 - val_loss:
1.6837
Epoch 11/40
28/28 [=====] - 102s 4s/step - loss: 0.1096 - val_loss:
1.3304
Epoch 12/40
28/28 [=====] - 103s 4s/step - loss: 0.1067 - val_loss:
0.6116
Epoch 13/40
28/28 [=====] - 103s 4s/step - loss: 0.0998 - val_loss:
0.5292
Epoch 14/40
28/28 [=====] - 106s 4s/step - loss: 0.1001 - val_loss:
0.1870
Epoch 15/40
28/28 [=====] - 104s 4s/step - loss: 0.0953 - val_loss:
0.1570
Epoch 16/40
28/28 [=====] - 105s 4s/step - loss: 0.0961 - val_loss:
0.1734
Epoch 17/40
28/28 [=====] - 100s 4s/step - loss: 0.0923 - val_loss:

```

```


0.1324
Epoch 18/40
28/28 [=====] - 98s 4s/step - loss: 0.0913 - val_loss:
0.1925
Epoch 19/40
28/28 [=====] - 99s 4s/step - loss: 0.0892 - val_loss:
0.1102
Epoch 20/40
28/28 [=====] - 98s 4s/step - loss: 0.0874 - val_loss:
0.2341
Epoch 21/40
28/28 [=====] - 98s 3s/step - loss: 0.0847 - val_loss:
0.1613
Epoch 22/40
28/28 [=====] - 98s 3s/step - loss: 0.0859 - val_loss:
0.1802
Epoch 23/40
28/28 [=====] - 98s 4s/step - loss: 0.0825 - val_loss:
0.1054
Epoch 24/40
28/28 [=====] - 101s 4s/step - loss: 0.0838 - val_loss:
0.1927
Epoch 25/40
28/28 [=====] - 101s 4s/step - loss: 0.0799 - val_loss:
0.0973
Epoch 26/40
28/28 [=====] - 99s 4s/step - loss: 0.0796 - val_loss:
0.1045
Epoch 27/40
28/28 [=====] - 102s 4s/step - loss: 0.0801 - val_loss:
0.0918
Epoch 28/40
28/28 [=====] - 99s 4s/step - loss: 0.0783 - val_loss:
0.1049
Epoch 29/40
28/28 [=====] - 98s 4s/step - loss: 0.0752 - val_loss:
0.0970
Epoch 30/40
28/28 [=====] - 98s 4s/step - loss: 0.0766 - val_loss:
0.0905
Epoch 31/40
28/28 [=====] - 97s 3s/step - loss: 0.0761 - val_loss:
0.1182
Epoch 32/40
28/28 [=====] - 103s 4s/step - loss: 0.0756 - val_loss:
0.0927
Epoch 33/40
28/28 [=====] - 97s 3s/step - loss: 0.0729 - val_loss:

```

```
0.1162
Epoch 34/40
28/28 [=====] - 98s 3s/step - loss: 0.0719 - val_loss:
0.0972
Epoch 35/40
28/28 [=====] - 98s 4s/step - loss: 0.0725 - val_loss:
0.1478
Epoch 36/40
28/28 [=====] - 99s 4s/step - loss: 0.0721 - val_loss:
0.0890
Epoch 37/40
28/28 [=====] - 98s 4s/step - loss: 0.0707 - val_loss:
0.1563
Epoch 38/40
28/28 [=====] - 98s 3s/step - loss: 0.0701 - val_loss:
0.1156
Epoch 39/40
28/28 [=====] - 98s 3s/step - loss: 0.0708 - val_loss:
0.0835
Epoch 40/40
28/28 [=====] - 98s 4s/step - loss: 0.0684 - val_loss:
0.2066
```

0.6 Plotting the Train and Cross-validation loss graph

```
[8]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



0.7 Comparing Groundtruth segmentation mask and Nural Networks's mask

```
[12]: val_gen = SeaKing(batch_size, img_size, val_input_img_paths,
    ↪ val_target_img_paths)
val_preds = model.predict(val_gen)

i = 11

display(Image(filename=val_input_img_paths[i]))

# Display ground-truth target mask
img = ImageOps.autocontrast(load_img(val_target_img_paths[i]))
display(img)

# Display mask predicted by our model
display_mask(i)
```

8/8 [=====] - 8s 1s/step



